

Setu

A multilingual AI platform that turns citizen infrastructure complaints into an explainable, ranked priority list for policymakers.

The problem

Governments across India struggle to consolidate citizen feedback and align it with national infrastructure priorities. Development requests live in fragmented systems — leading to misaligned public spending, unaddressed infrastructure gaps, and no reliable way to measure the impact of digital public infrastructure initiatives.

Fragmented

Complaints scattered across calls, forms, offices — no single view

Multilingual

Citizens write in Hindi, English, and mixed regional languages


Unmeasured

No way to tell if public spending is reaching real need

What Setu does

- 1 Ingests** Citizen complaints via text — voice architected for — in any Indian language, through an open API
- 2 Understands** Gemini extracts topic, district, urgency — refusing to guess when information isn't stated
- 3 Fuses** Complaints joined with real census data — population, infrastructure access — by district
- 4 Surfaces** A ranked, explainable priority list a policymaker can act on — and interrogate

Architecture



Write before process

Every complaint is durably stored before anything else happens — nothing is lost if a step fails

Micro-batching, not one call per complaint

Gemini's real limit is 5 requests/minute — batching turns that into ~50 complaints/minute

A harness, not blind trust

Every Gemini response is validated — IDs matched, enums checked, districts confirmed real

LIVE DEMO

Submit a complaint. See it ranked with
a reason.

setu-alpha-eight.vercel.app

Not just working — tested

Disambiguation, caught and fixed

A streetlight complaint was misclassified as "electricity." We added explicit rules to the prompt — verified fixed on the next test.

Scoring, rebalanced

A small district with 11 complaints outranked one with 28 by 10×. Log-compressing both terms brought the gap to a defensible 2.7×.

Under-reporting, corrected for

Districts with low internet and literacy access are structurally less able to complain digitally. Their scores are weighted up, not down.

Duplicates, collapsed

14 reports in one district were grouped into 12 distinct issues before scoring — so repetition doesn't inflate priority.

Honest about scope

Census 2011

India's most recent complete census. Population scaled by a sourced 2011→2026 growth multiplier, stated explicitly as a national average, not fabricated precision.

Demo data

Synthetic complaints, clearly labeled, run through the exact same live pipeline real citizen submissions would use — no shortcut around ingestion or extraction.

Not yet built

Voice ingestion (architected for Cloud Speech-to-Text), a large statistical evaluation set, personal data redaction.

Tech stack

Layer	Choice
API	FastAPI
Database	PostgreSQL, Alembic migrations
AI / NLP	Gemini 3.8 Flash, batched extraction
Public data	Census 2011, district-level indicators
Deployment	Vercel (dashboard) · Render (API) · Neon (Postgres)

DESIGNED AS A DIGITAL PUBLIC GOOD

An open ingestion layer any channel can plug into — a WhatsApp bot, an IVR system, a web form — not a single closed app.

Open API, not a walled garden

Explainable scoring, not a black box

Built to scale by adding workers, not redesigning

Setu

Prototype

setu-alpha-eight.vercel.app

Repository

github.com/AnasFaaiz/setu

Track

AI for Digital Public Infrastructure & Governance